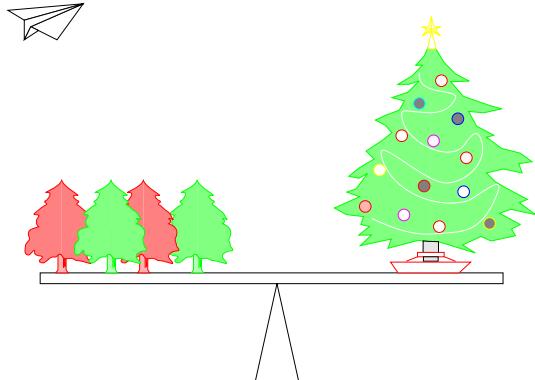


2-3-4 Trees and Red-Black Trees



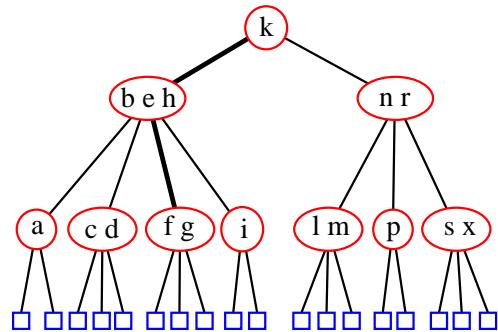
erm

204



2-3-4 Trees Revealed

- Nodes store 1, 2, or 3 keys and have 2, 3, or 4 children, respectively
- All leaves have the same depth



$$\frac{1}{2} \log(N+1) \leq \text{height} \leq \log(N+1)$$

erm

205

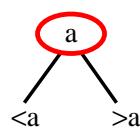


2-3-4 Tree Nodes

- Introduction of nodes with more than 1 key, and more than 2 children

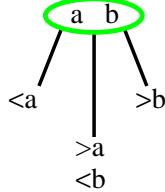
2-node:

- same as a binary node



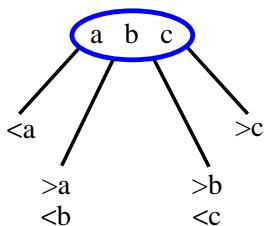
3-node:

- 2 keys, 3 links



4-node:

- 3 keys, 4 links



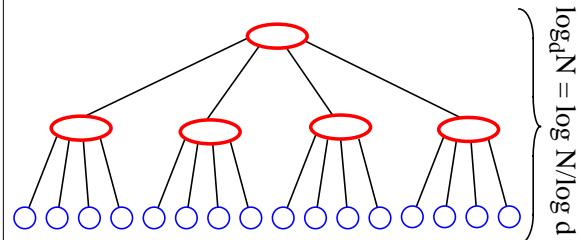
erm

206



Why 2-3-4?

- Why not minimize height by maximizing children in a “d-tree”?
- Let each node have d children so that we get $O(\log_d N)$ search time! Right?



- That means if $d = N^{1/2}$, we get a height of 2
- However, searching out the correct child on each level requires $O(\log N^{1/2})$ by binary search
- $2 \log N^{1/2} = O(\log N)$ which is not as good as we had hoped for!
- 2-3-4-trees will guarantee $O(\log N)$ height using only 2, 3, or 4 children per node

erm

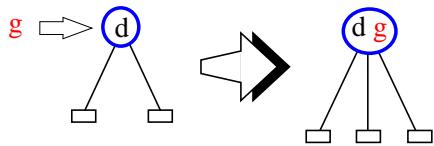
207



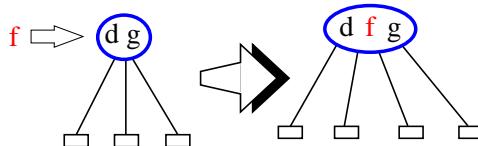
Insertion into 2-3-4 Trees

- Insert the **new key** at the **lowest internal node reached** in the search

- 2-node** becomes **3-node**



- 3-node** becomes **4-node**

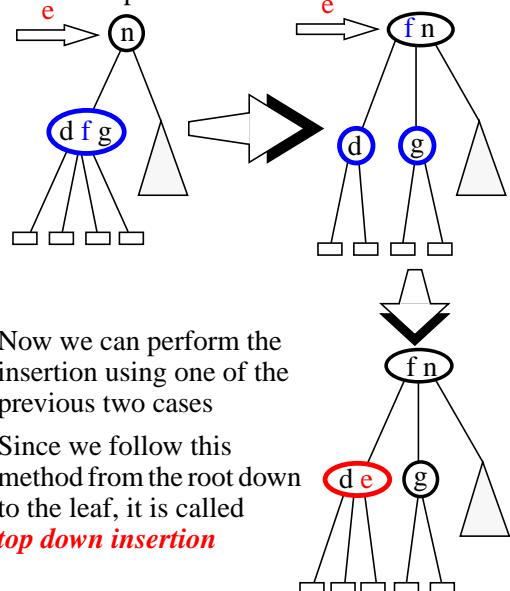


- What about a **4-node**?

- We can't insert another key!

Top Down Insertion

- In our way down the tree, whenever we reach a **4-node**, we **break it up** into two **2-nodes**, and move the middle element up into the parent node

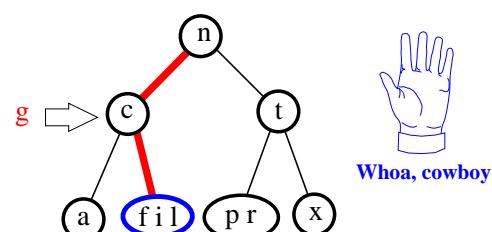
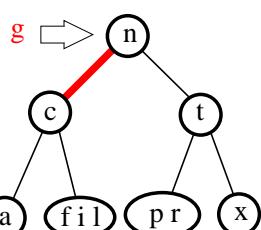
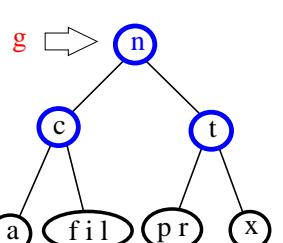
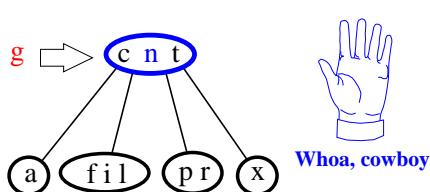


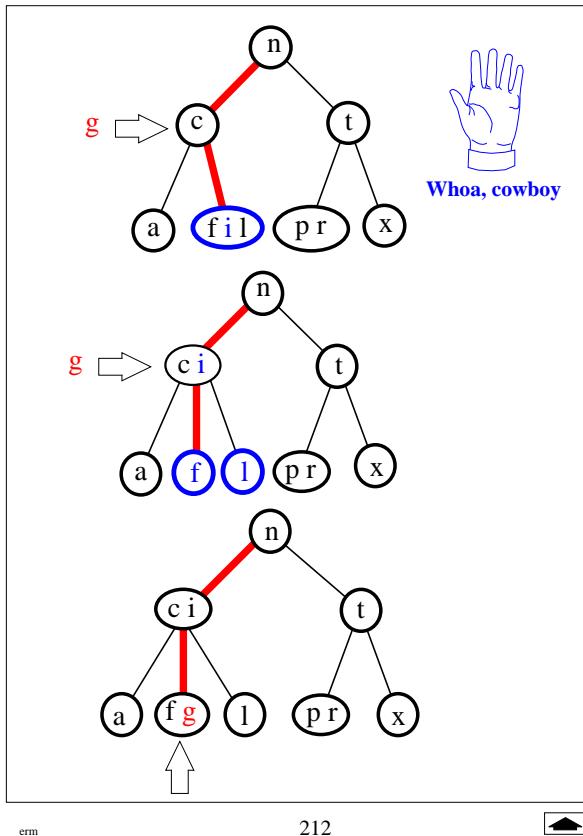
- Now we can perform the insertion using one of the previous two cases

- Since we follow this method from the root down to the leaf, it is called **top down insertion**

Splitting the Tree

As we travel down the tree, if we encounter any **4-node** we will break it up into **2-nodes**. This guarantees that we will never have the problem of inserting the middle element of a former **4-node** into its parent **4-node**.





212



erm

213



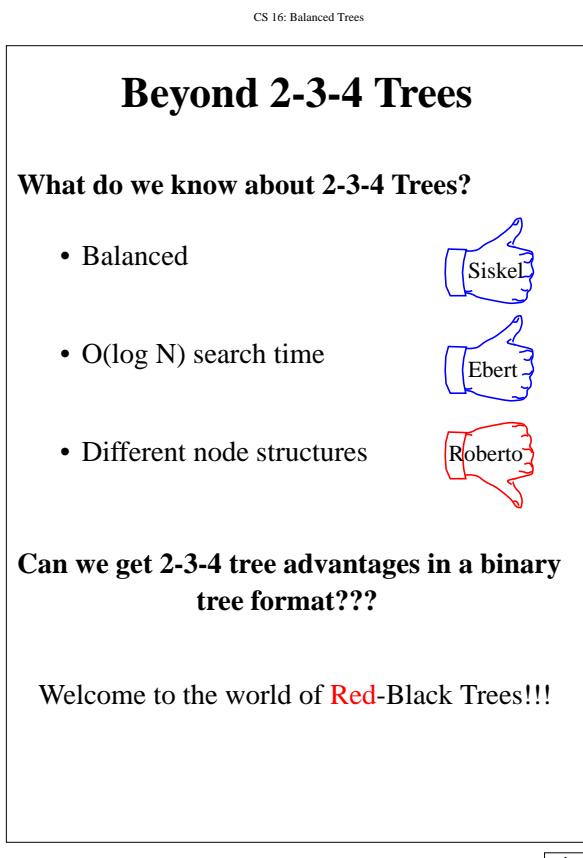
Time Complexity of Insertion in 2-3-4 Trees

Time complexity:

- A search visits $O(\log N)$ nodes
- An insertion requires $O(\log N)$ node splits
- Each node split takes constant time
- Hence, operations **Search** and **Insert** each take time $O(\log N)$

Notes:

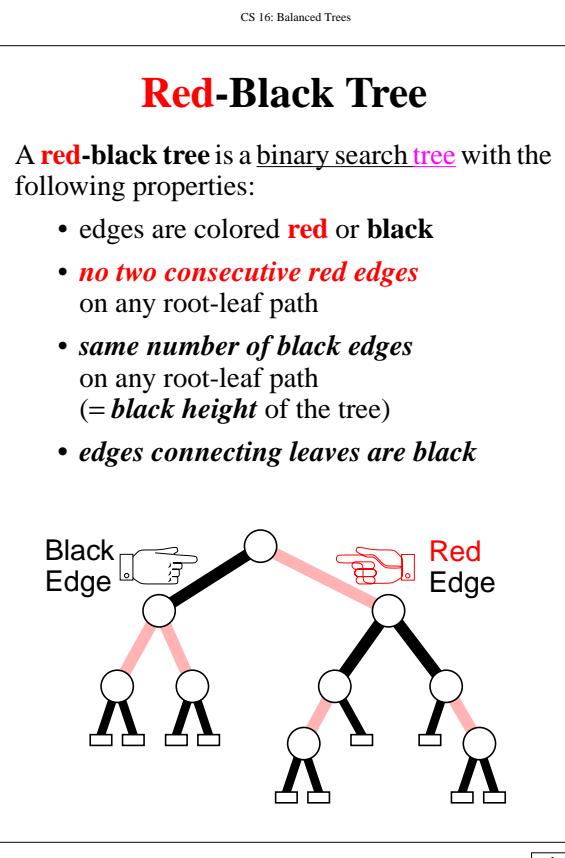
- Instead of doing splits top-down, we can perform them bottom-up starting at the insertion node, and only when needed. This is called **bottom-up** insertion.
- A deletion can be performed by *fusing* nodes (inverse of splitting), and takes $O(\log N)$ time



214



erm



215

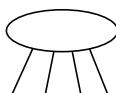


erm

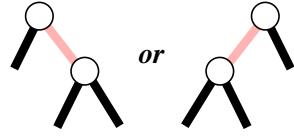
2-3-4 Tree Evolution

Note how 2-3-4 trees relate to **red-black trees**

2-3-4



Red-Black



Now we see red-black trees are just a way of representing 2-3-4 trees!

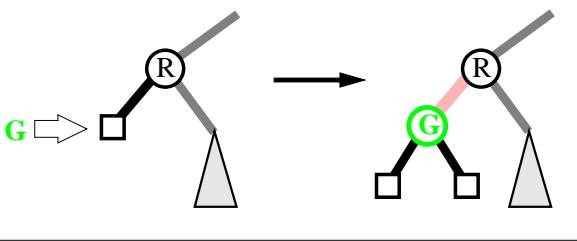
erm

216



Insertion into Red-Black Trees

1. Perform a standard search to find the leaf where the key should be added
2. Replace the leaf with an internal node with the new key
3. Color the incoming edge of the new node **red**
4. Add two new leaves, and color their incoming edges **black**
5. If the parent had an incoming **red** edge, we now have two consecutive **red** edges! We must reorganize tree to remove that violation. What must be done depends on the sibling of the parent.



erm

218



More Red-Black Tree Properties

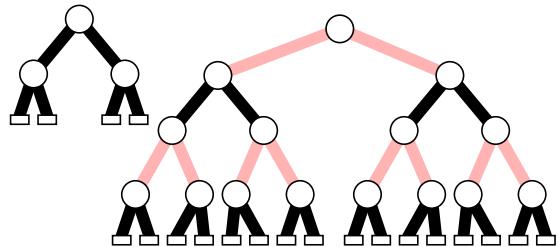
N # of internal nodes

L # leaves (= N + 1)

H height

B black height

Property 1: $2^B \leq N + 1 \leq 4^B$



Property 2: $\frac{1}{2}\log(N + 1) \leq B \leq \log(N + 1)$

Property 3: $\log(N + 1) \leq H \leq 2\log(N + 1)$

This implies that searches take time **O(logN)**!

erm

217



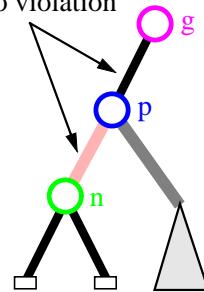
Insertion - Plain and Simple

Let:

- **n** be the new node
- **p** be its parent
- **g** be its grandparent

Case 1: Incoming edge of **p is black**

No violation



STOP!

Pretty easy, huh?

Well... it gets messier...

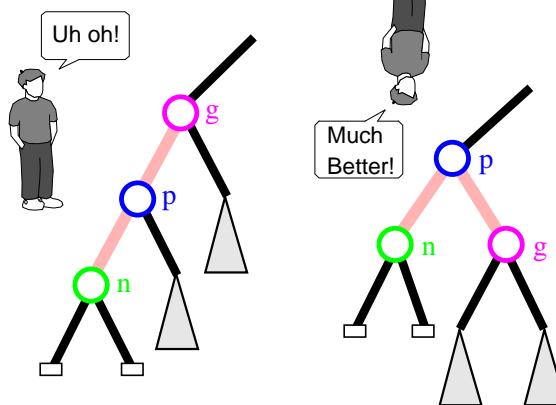
erm

219



Restructuring

Case 2: Incoming edge of **p** is **red**, and its sibling is **black**



We call this a “*right rotation*”

- No further work on tree is necessary
- Inorder remains unchanged
- Tree becomes more balanced
- No two consecutive **red** edges!

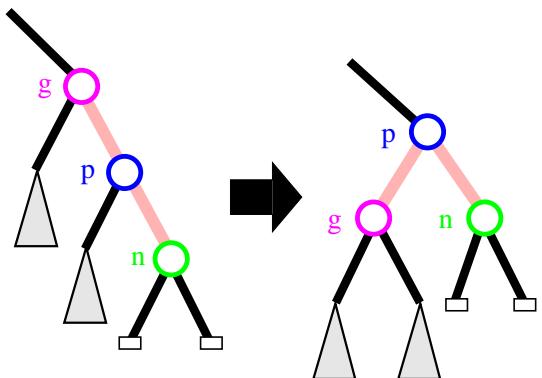
erm

220



More Rotations

Similar to a right rotation, we can do a “**left rotation**”...



Simple, huh?

erm

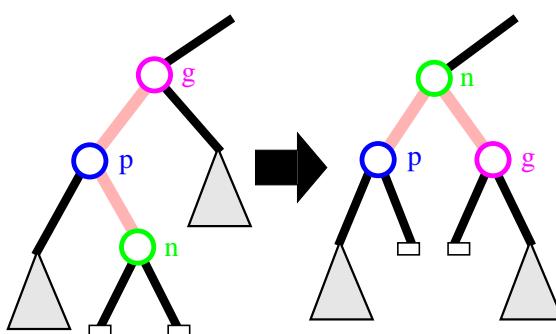
221



Double Rotations

What if the new node is between its parent and grandparent in the inorder sequence?

We must perform a “**double rotation**”
(which is no more difficult than a “**single**” one)



This would be called a
“left-right double rotation”

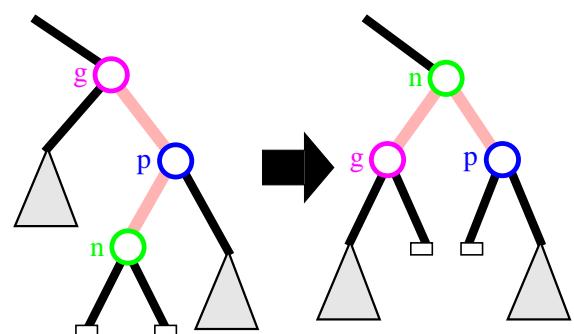
erm

222



Last of the Rotations

And this would be called a
“right-left double rotation”



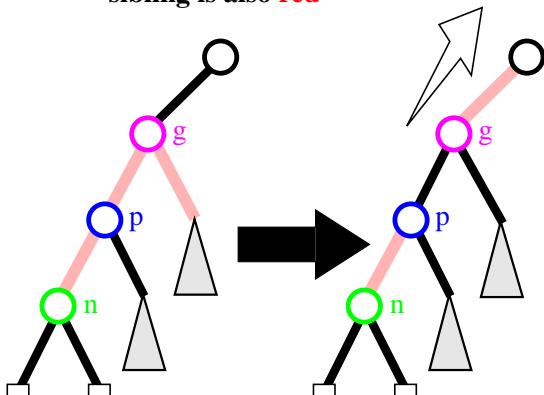
erm

223



Bottom-Up Rebalancing

Case 3: Incoming edge of **p** is **red** and its sibling is also **red**



- We call this a “**promotion**”
- Note how the black depth remains unchanged for all of the descendants of **g**
- This process will continue upward beyond **g** if necessary: rename **g** as **n** and repeat.

224



erm

225

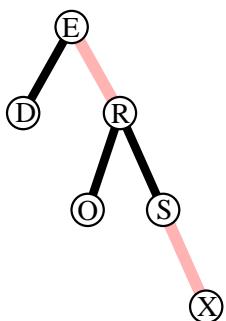


Summary of Insertion

- If **two red edges** are present, we do either
 - a **restructuring** (with a simple or double rotation) and **stop**, or
 - a **promotion** and **continue**
- A **restructuring** takes **constant time** and is performed at most once. It reorganizes an off-balanced section of the tree.
- **Promotions** may continue up the tree and are executed **$O(\log N)$** times.
- The **time complexity** of an insertion is **$O(\log N)$** .

An Example

Start by inserting “REDSOX” into an empty tree



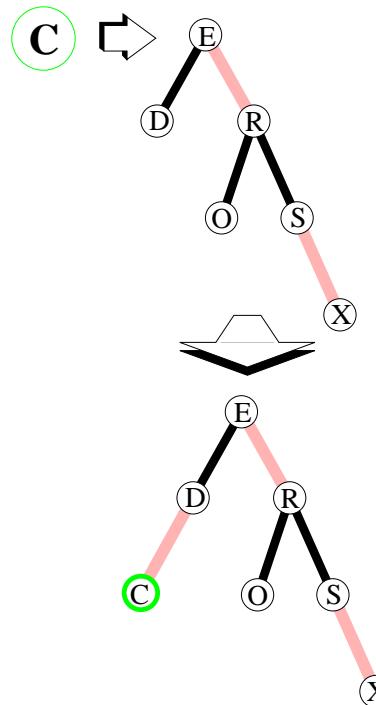
Now, let's insert “C U B S”...

226



erm

A Cool Example

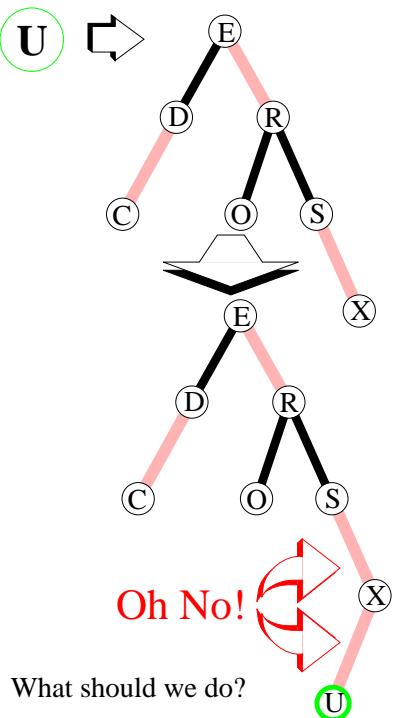


227



erm

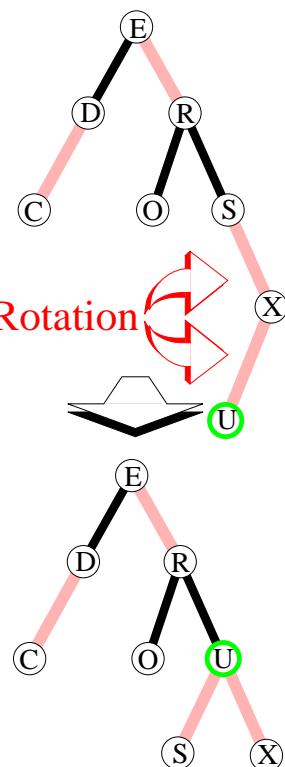
An Unbelievable Example



228



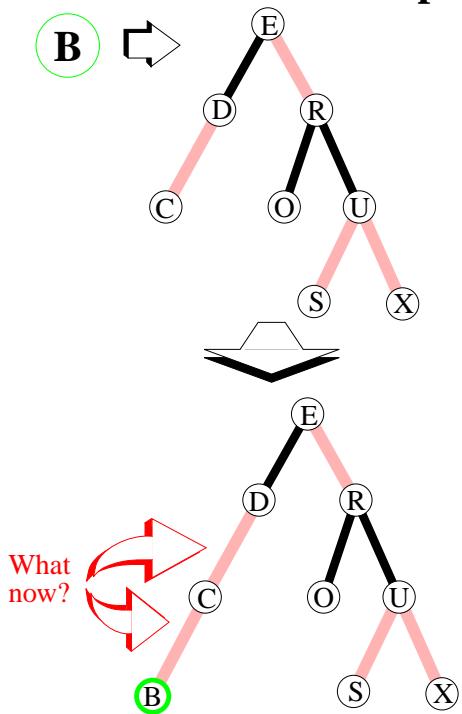
Double Rotation



229



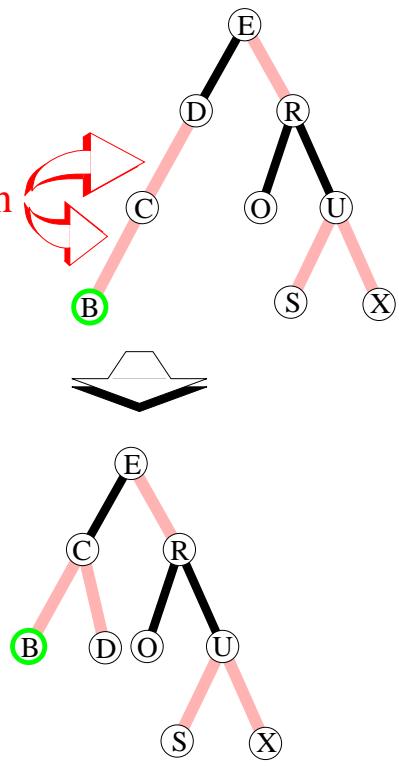
A Beautiful Example



230



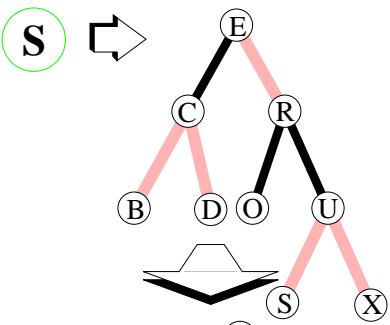
Rotation



231

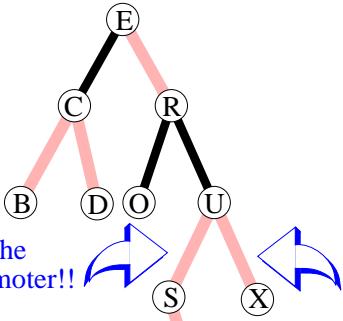


A Super Example



Holy Consecutive Red Edges, Batman!  We could've placed it on either side

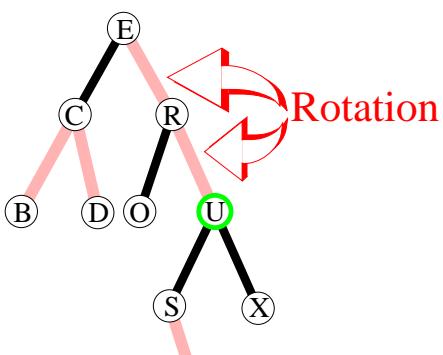
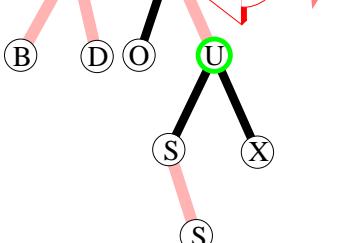
We could've placed it on either side



Us
Bat-P



BIFF!



The SUN lab
and Red-*Bat*
trees are safe...
...for now!!!

