

Hashing

What is it?

A form of narcotic intake?

A side order for your eggs?

A combination of the two?

Another Solution

- We can do better, with a *Hashtable* -- O(1) expected time, O(N+M) space, where M is table size
- Like an array, but come up with a function to map the large range into one which we can manage
 - e.g., take the original key, modulo the (relatively small) size of the array, and use that as an index
- Insert (863-7639, Roberto) into a hashed array with, say, five slots
 - $8637639 \bmod 5 = 4$, so (863-7639, Roberto) goes in slot 4 of the hash table

(null)	(null)	(null)	(null)	Roberto
0	1	2	3	4


- A lookup uses the same process: hash the query key, then check the array at that slot
- Insert (863-9350, Gordon)
- And insert (863-2234, Gordon). Don't skip this example!

Problem

- RT&T is a large phone company, and they want to provide caller ID capability:
 - given a phone number, return the caller's name
 - phone numbers are in the range R=0 to $10^7 - 1$
 - want to do this as efficiently as possible (\$\$\$)
- A few suboptimal ways to design this dictionary:
 - an array indexed by key: takes O(1) time, O(N+R) space -- huge amount of wasted space

(null)	(null)	...	Roberto	...	(null)
000-0000	000-0001	...	863-7639	...	999-9999

- a linked list: takes O(N) time, O(N) space




- a balanced binary tree: O(lg N) time, O(N) space
(you want fancy pictures here too? so read the slides from the RedBlack help session).

Collision Resolution

- How to deal with two keys which hash to the same spot in the array?
- Use *chaining*

- Set up an array of links (a **table**), indexed by the keys, to **lists** of items with the same key



- Most efficient (time-wise) collision resolution
 - we'll talk about others later which use less space

Pseudo-code

- Any dictionary has 3 basic methods, and the constructor:
`init`
`insert`
`find`
`remove`
- Init
`create table of M lists`
- Insert(K)
`index = h(K)`
`insert into table[index]`
- Find(K)
`index = h(K)`
walk down list at `table[index]`, looking for a match
return what was found (or error)
- Remove(K)
`index = h(K)`
walk down list at `table[index]`, looking for a match
remove what was found (or error)

Hash Functions


- Need to choose a good hash function
 - quick to compute
 - distributes keys uniformly throughout the table
- How to deal with hashing non-integer keys:
 - find some way of turning the keys into integers
 - in our example, remove the hyphen in 863-7639 to get 8637639!
 - for a string, add up the ASCII values of the characters of your string
 - then use a standard hash function on the integers
- Use the remainder
 - $h(K) = K \text{ mod } M$
 - K is the key, M the size of the table
- Need to choose M
- **M = b^e (bad)**
 - if M is a power of two, $h(K)$ gives the e least significant bits of K
 - all keys with the same ending go to the same place
- **M prime (good)**
 - helps ensure uniform distribution
 - take a number theory class to understand why

Hash Functions (cont.)

- Mid-Square
 - $h(K) = \text{middle digits of } K^2$
- I.E. Table size power of 10
 - $h(4150130) = 21526 \mathbf{4436} 17100$
 - $h(415013034) = 526447 \mathbf{3522} 151420$
 - $h(1150130) = 13454 \mathbf{2361} 7100$
- I.E. Table power is power of 2
 - $h(1001) = 10 \mathbf{100} 01$
 - $h(1011) = 11 \mathbf{110} 01$
 - $h(1101) = 101 \mathbf{010} 01$

More on Collisions

- A key is mapped to an already occupied table location
 - what to do?!?
- Use a collision handling technique
- We've seen *Chaining*
- Can also use *Open Addressing*
 - Double Hashing
 - Linear Probing



Linear Probing

- If the current location is used, try the next table location

```
linear_probing_insert(K)
  if (table is full) error

  probe = h(K)

  while (table[probe] occupied)
    probe = (probe + 1) mod M


  table[probe] = K
```

- Lookups walk along table until the key or an empty slot is found
- Uses less memory than chaining
 - don't have to store all those links
- Slower than chaining
 - may have to walk along table for a long way
- A real pain to delete from
 - either mark the deleted slot
 - or fill in the slot by shifting some elements down

Hashing

9

18 41 22 44 59 32 31 73



Hashing

10

Double Hashing

- Use two hash functions
- If M is prime, eventually will examine every position in the table

```
double_hash_insert(K)
  if(table is full) error

  probe = h1(K)
  offset = h2(K)

  while (table[probe] occupied)
    probe = (probe + offset) mod M

  table[probe] = K
```

- Many of same (dis)advantages as linear probing
- Distributes keys more uniformly than linear probing does


Hashing

11

Double Hashing Example

- $h1(K) = K \bmod 13$
- $h2(K) = 8 - K \bmod 8$
 - we want $h2$ to be an offset to add

18 41 22 44 59 32 31 73



Hashing

12

Theoretical Results

- Let $\alpha = N/M$
 - the load factor: average number of keys per array index
- Analysis is probabilistic, rather than worst-case

Expected Number of Probes


	<i>not found</i>	<i>found</i>
Chaining	$1 + \alpha$	$1 + \frac{\alpha}{2}$
Linear Probing	$\frac{1}{2} + \frac{1}{2(1-\alpha)^2}$	$\frac{1}{2} + \frac{1}{2(1-\alpha)}$
Double Hashing	$\frac{1}{(1-\alpha)}$	$\frac{1}{\alpha} \ln \frac{1}{1-\alpha}$

Hashing

13

Pretty Graph

Expected Number of Probes vs. Load Factor



Hashing

14